

---

## Labsheet - 2

### Linear Regression and Polynomial Curve Fitting

Machine Learning  
BITS F464  
I Semester 2021-22

---

In this Lab-sheet, we will work on Linear Regression model and then we will see basics about polynomial regression.

### Linear Regression:

The general mathematical equation for a linear regression model is –

$$Y = ax + b$$

here,

y is the response variable.

x is the predictor variable.

a,b are weights associated with input.

Let's take a simple example of sine wave for which `lm()` function is used in R.

### Input Data:

Generate some data by using the function `sin(2πx)` in `[0,1]`

```
x = seq(0,1,length=11)
y = sin(2*pi*x)
Some noise is generated and added to the real signal (y):
noise = y + rnorm(11, sd=0.3)
#we will make y the response variable and x the predictor
#the response variable is usually on the y-axis
plot(x,noise,pch=19)
```

The basic syntax for `lm()` function in linear regression is -

`lm(formula,data)`

```
#fit first degree polynomial equation:
lm1 <- lm(y~x)
#second degree
```

```
lm2 <- lm(y~poly(x,2))
```

Now try the model it with degree varying from 3 to 9 and plot the fits.

**Write down your observations as degree increases from 1 to 9.**

Now increase the number of training points, N as mentioned below for M=9

N=10

N=15

N=30

N=50

N=100

**Write down your observations as N increases.**

## Overfitting Problem:

### 1) Plot a graph of training error and testing error

```
#getting the test and training errors
lms = list(lm1,lm2,lm3,lm4,lm5,lm6,lm7,lm8,lm9)
# set up two null vectors for the errors
train = test = rep(0,10) # train = vector(length=10) also works
# this is SSE/10
for(i in 1:10){train[i]=var(lms[i][[1]]$residuals)}
# now compute similarly scaled prediction error for test data
newy = sin(2*pi*x)+rnorm(11,sd=.3) # note we used the same x's
for(j in 1:10){test[j] = (1/10)*sum((newy-predict(lms[j]
[[1]],newdata=data.frame(x)))^2)}
lines (1:10,train, ylim=c(0,.5), col='black',lwd=3)
lines(1:10,test,pch="X", col='red',lwd=3)
cbind(train,test)
```

### 2) Show all coefficients of the polynomial in a tabular form. And write down your observations.

Coefficients are the intercept and the slope of the regression line, but more informative results about the model can find by:

```
#summary
summary(lm1)
#coefficient
C1 <- coef(lm1)
```

Use data frame for table generation

```
max_length <- max(c(length(C1), length(C2))) # Find out maximum
length for unequal table
```

```
#Now, we can use this information to create a data frame with NA
values at the bottom of each column that is shorter as the
maximum length.
data <- data.frame(col1 = c(C1,rep(NA, max_length - length(C1))),
                   col2 = c(C2,rep(NA, max_length - length(C2))))
```

## Regularization:

Linear regression algorithm works by selecting coefficients for each independent variable that minimizes a loss function. However, if the coefficients are large, they can lead to over-fitting on the training dataset, and such a model will not generalize well on the unseen test data. To overcome this shortcoming, we'll do regularization, which penalizes large coefficients.

We will be using the `glmnet()` package to build the regularized regression models. The `glmnet` function does not work with dataframes, so we need to create a numeric matrix for the training features and a vector of target values.

```
lambdas <- 10^seq(3, -2, by = -.1)
fit <- glmnet(x, y, alpha = 0, lambda = lambdas)
```

## Exercise

- Load “mtcar” dataset and split it into training and testing part. Plot a graph with various polynomial degree on training dataset. Then choose suitable degree for testing data.
  - Apply ridge regularization by using `glmnet` function and plot the graph between `lambda` & error.
-